CS324: Deep Learning
Assignment 1

Jianguo Zhang

March 6, 2024

1 Part I: the perceptron (25 points)

In this first task you're asked to implement and test a simple artificial neuron: a perceptron (see **perceptron_tutorial.pdf**).

1.1 Task 1

Generate a dataset of points in $\mathbb{R}^2$. To do this, define two Gaussian distributions and sample 100 points from each. Your dataset should then contain a total of 200 points, 100 from each distribution. Keep 80 points per distribution as the training (160 in total), 20 for the test (40 in total).

1.2 Task 2

Implement the perceptron following the specs in **perceptron.py** and the stantard algorithm section in **perceptron_tutorial.pdf**.

1.3 Task 3

Train the perceptron on the training data (160 points) and test in on the remaining 40 test points. Compute the classification accuracy on the test set.

1.4 Task 4

Experiment with different sets of points (generated as described in Task 1). What happens during the training if the means of the two Gaussians are too close and/or if their variance is too high?

2 Part II: the mutli-layer perceptron (60 points)

In this second part of Assignment I you're asked to implement a multi-layer perceptron using numpy. Using scikit-learn and the `make_moons` method[1], create a dataset of $1,000$ two-dimensional points. Let S denote the dataset, i.e., the set of tuples $\{(x^{(0),s}, t^s)\}_{s=1}^{S}$, where $x^{(0),s}$ is the s-th element of the dataset and t^s is its label. Further let d_0 be the dimension of the input space and d_n the dimension of the output space. In this assignment we want the labels to be one-hot encoded[2]. The network you will build will have N layers (including the output layer). In particular, the structure will be as follows:

---

[1] https://scikit-learn.org/stable/modules/generated/sklearn.datasets.make_moons.html#sklearn.datasets.make_moons
[2] Remember to transform the original dataset labels using one-hot encoding https://en.wikipedia.org/wiki/One-hot

- Each layer $l = 1, \cdots, N$ first applies the affine mapping

$$\tilde{x}^{(l)} = W^{(l)} x^{(l-1)} + b^{(l)},$$

where $W^{(l)} \in \mathbb{R}^{d_l \times d_{(l-1)}}$ is the matrix of the weight parameters and $b^{(l)} \in \mathbb{R}^{d_l}$ is the vector of biases. Given $\tilde{x}^{(l)}$, the activation of the l-th layer is computed using a ReLU unit

$$x^{(l)} = \max(0, \tilde{x}^{(l)}).$$

- The output layer (i.e., the N-th layer) first applies the affine mapping

$$\tilde{x}^{(N)} = W^{(N)} x^{(N-1)} + b^{(N)},$$

and then uses the softmax activation function (instead of the ReLU of the previous layers) to compute a valid probability mass function (pmf)

$$x^{(N)} = \text{softmax}(\tilde{x}^{(N)}) = \frac{\exp(\tilde{x}^{(N)})}{\sum_{i=1}^{d_N} \exp(\tilde{x}^{(N)})_i}.$$

Note that both max and exp are element-wise operations.

- Finally, compute the cross entropy loss L between the predicted and the actual label,

$$L(x^{(N)}, t) = -\sum_i t_i \log x_i^{(N)}.$$

2.1 Task 1

Implement the MLP architecture by completing the files **mlp_numpy.py** and **modules.py**.

2.2 Task 2

Implement training and testing script in **train_mlp_numpy.py**. (Please keep 80% of the dataset for training and the remaining 20% for testing. Note that this is a random split of 80% and 20%)

2.3 Task 3

Using the default values of the parameters, report the results of your experiments using a jupyter notebook where you show the accuracy curves for both training and test data.

3 Part III: stochastic gradient descent (15 points)

In this third part of Assignment I you will implement an alternative traning method in **train_mlp_numpy.py** based on stochastic gradient descent.

3.1 Task 1

Modify the train method in **train_mlp_numpy.py** to accept a parameter that allows the user to specify if the training has to be performed using batch gradient descent (which you should have implemented in Part II) or stochastic gradient descent (batch size equals 1).

3.2 Task 2

Using the default values of the parameters, report the results of your experiments using a jupyter notebook where you show the accuracy curves for both training and test data. In your report, you should also show and analyze the influence of batch size from 1 to a relative large number.

4 Grading Rule

- The score ratio for report and code is 4:6.

- The report should not be less than 4 pages and pictures in the report should not be more than 1/3 of the report. You can put the extra pictures and tables in the appendix.

- The report should include but not limited to:
    - Theoretical analysis of learning process of the perceptron.
    - Theoretical analysis of forward and backward propagation of each layer in multi-layer perceptron.
    - Loss curve of training and testing.
    - Analysis of results you got.

- The report will be scored according to the quality of the analysis.

5 Submission Instructions

The submission will include:

- A written report describing what you did, the results and your analysis.

- Code for producing **all** results for all parts and tasks.

- Instructions on how to run the code.

Create a ZIP archive with the submission of Assignment 1 (all parts and tasks). Give the ZIP file the name **studentnumber_assignment1.zip**, where you insert your student number. Please submit the archive through the Blackboard.

Make sure all files needed to run your code are included or you may be given 0 points for it.

The deadline for assignment 1 (all parts and all tasks) is the 28th of March 2024 at 23:55 (Beijing Time).